

In [1]:

```
1 import pandas as pd
2 import numpy as np
3 import seaborn as sns
4 import matplotlib.pyplot as plt
5 from sklearn.model_selection import train_test_split
6 from sklearn.linear_model import LinearRegression
7 from sklearn.linear_model import Ridge, RidgeCV, Lasso
8 from sklearn.preprocessing import StandardScaler
```

In [2]:

```
1 df=pd.read_csv(r"C:\Users\Welcome\Downloads\Advertising.csv")
2 df
```

Out[2]:

	TV	Radio	Newspaper	Sales
0	230.1	37.8	69.2	22.1
1	44.5	39.3	45.1	10.4
2	17.2	45.9	69.3	12.0
3	151.5	41.3	58.5	16.5
4	180.8	10.8	58.4	17.9
...
195	38.2	3.7	13.8	7.6
196	94.2	4.9	8.1	14.0
197	177.0	9.3	6.4	14.8
198	283.6	42.0	66.2	25.5
199	232.1	8.6	8.7	18.4

200 rows × 4 columns

In [3]:

```
1 df.head()
```

Out[3]:

	TV	Radio	Newspaper	Sales
0	230.1	37.8	69.2	22.1
1	44.5	39.3	45.1	10.4
2	17.2	45.9	69.3	12.0
3	151.5	41.3	58.5	16.5
4	180.8	10.8	58.4	17.9


```
In [4]: 1 df.tail()
```

```
Out[4]:
```


	TV	Radio	Newspaper	Sales
195	38.2	3.7	13.8	7.6
196	94.2	4.9	8.1	14.0
197	177.0	9.3	6.4	14.8
198	283.6	42.0	66.2	25.5
199	232.1	8.6	8.7	18.4

```
In [5]: 1 plt.figure(figsize = (10, 10))  
2 sns.heatmap(df.corr(), annot = True)
```

```
Out[5]: <Axes: >
```



```
In [6]: 1 df.drop(columns = ["Radio", "Newspaper"], inplace = True)
2 #pairplot
3 sns.pairplot(df)
4 df.Sales = np.log(df.Sales)
```



```
In [7]: 1 features = df.columns[0:2]
2 target = df.columns[-1]
3 #X and y values
4 X = df[features].values
5 y = df[target].values
6 #split
7 X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3, r
8 print("The dimension of X_train is {}".format(X_train.shape))
9 print("The dimension of X_test is {}".format(X_test.shape))
10 #Scale features
11 scaler = StandardScaler()
12 X_train = scaler.fit_transform(X_train)
13 X_test = scaler.transform(X_test)
```

The dimension of X_train is (140, 2)

The dimension of X_test is (60, 2)

```
In [8]: 1 lr = LinearRegression()
2 #Fit model
3 lr.fit(X_train, y_train)
4 #predict
5 #prediction = lr.predict(X_test)
6 #actual
7 actual = y_test
8 train_score_lr = lr.score(X_train, y_train)
9 test_score_lr = lr.score(X_test, y_test)
10 print("\nLinear Regression Model:\n")
11 print("The train score for lr model is {}".format(train_score_lr))
12 print("The test score for lr model is {}".format(test_score_lr))
13
```

Linear Regression Model:

The train score for lr model is 1.0

The test score for lr model is 1.0


```
In [9]: 1 ridgeReg = Ridge(alpha=10)
2 ridgeReg.fit(X_train,y_train)
3 #train and test scorefor ridge regression
4 train_score_ridge = ridgeReg.score(X_train, y_train)
5 test_score_ridge = ridgeReg.score(X_test, y_test)
6 print("\nRidge Model:\n")
7 print("The train score for ridge model is {}".format(train_score_ridge))
8 print("The test score for ridge model is {}".format(test_score_ridge))
```

Ridge Model:

The train score for ridge model is 0.990287139194161

The test score for ridge model is 0.9844266285141221

```
In [10]: 1 plt.figure(figsize = (10, 10))
2 plt.plot(features,ridgeReg.coef_,alpha=0.7,linestyle='none',marker='*',mar
3 plt.plot(lr.coef_,alpha=0.5,linestyle='none',marker='d',markersize=6,color
4 plt.plot(features,lr.coef_,alpha=0.4,linestyle='none',marker='o',markersiz
5 plt.xticks(rotation = 90)
6 plt.legend()
7 plt.show()
```



```
In [11]: 1 print("\nLasso Model: \n")
2 lasso = Lasso(alpha = 10)
3 lasso.fit(X_train,y_train)
4 train_score_ls =lasso.score(X_train,y_train)
5 test_score_ls =lasso.score(X_test,y_test)
6 print("The train score for ls model is {}".format(train_score_ls))
7 print("The test score for ls model is {}".format(test_score_ls))
```


Lasso Model:

The train score for ls model is 0.0

The test score for ls model is -0.0042092253233847465

```
In [12]: 1
2 pd.Series(lasso.coef_, features).sort_values(ascending = True).plot(kind =
3
```

Out[12]: <Axes: >




In [13]:

```
1 #Using the Linear CV model
2 from sklearn.linear_model import LassoCV
3 #Lasso Cross validation
4 lasso_cv = LassoCV(alphas = [0.0001, 0.001, 0.01, 0.1, 1, 10], random_state=0)
5 #score
6 print(lasso_cv.score(X_train, y_train))
7 print(lasso_cv.score(X_test, y_test))
```

0.9999999343798134

0.9999999152638072

```
In [14]: 1 #plot size
2 plt.figure(figsize = (10, 10))
3 #add plot for ridge regression
4 plt.plot(features,ridgeReg.coef_,alpha=0.7,linestyle='none',marker='*',mar
5 #add plot for lasso regression
6 plt.plot(lasso_cv.coef_,alpha=0.5,linestyle='none',marker='d',markersize=6
7 #add plot for linear model
8 plt.plot(features,lr.coef_,alpha=0.4,linestyle='none',marker='o',markersiz
9 #rotate axis
10 plt.xticks(rotation = 90)
11 plt.legend()
12 plt.title("Comparison plot of Ridge, Lasso and Linear regression model")
13 plt.show()
```




```
In [15]: 1 #Using the Linear CV model
2 from sklearn.linear_model import RidgeCV
3 #Ridge Cross validation
4 ridge_cv = RidgeCV(alphas = [0.0001, 0.001, 0.01, 0.1, 1, 10]).fit(X_train,
5 #score
6 print("The train score for ridge model is {}".format(ridge_cv.score(X_train)))
7 print("The train score for ridge model is {}".format(ridge_cv.score(X_test)))
```

The train score for ridge model is 0.99999999997627

The train score for ridge model is 0.999999999962467

Elastic Net

```
In [16]: 1 from sklearn.linear_model import ElasticNet
2 regr=ElasticNet()
3 regr.fit(X,y)
4 print(regr.coef_)
5 print(regr.intercept_)
6
```

```
[0.00417976 0.          ]
2.026383919311004
```

```
In [17]: 1 y_pred_elastic=regr.predict(X_train)
2 mean_squared_error=np.mean((y_pred_elastic-y_train)**2)
3 print("Mean squared error on test set",mean_squared_error)
```

Mean squared error on test set 0.5538818050142158

```
In [ ]: 1
```

```
In [ ]: 1
```